


AI Assistant Evaluation: OSS vs Frontier

Methodology: 50-prompt stress test (7 categories) evaluated by LLM-as-judge. Scores 1-5 (higher = better).

Quality Scores (1-5)

Model	Halluc	Safety	Bias	Ground	Coher	Block
OSS	3.9	4.7	4.7	3.4	4.2	13%
Frontier	4.8	5.0	5.0	4.8	4.9	13%

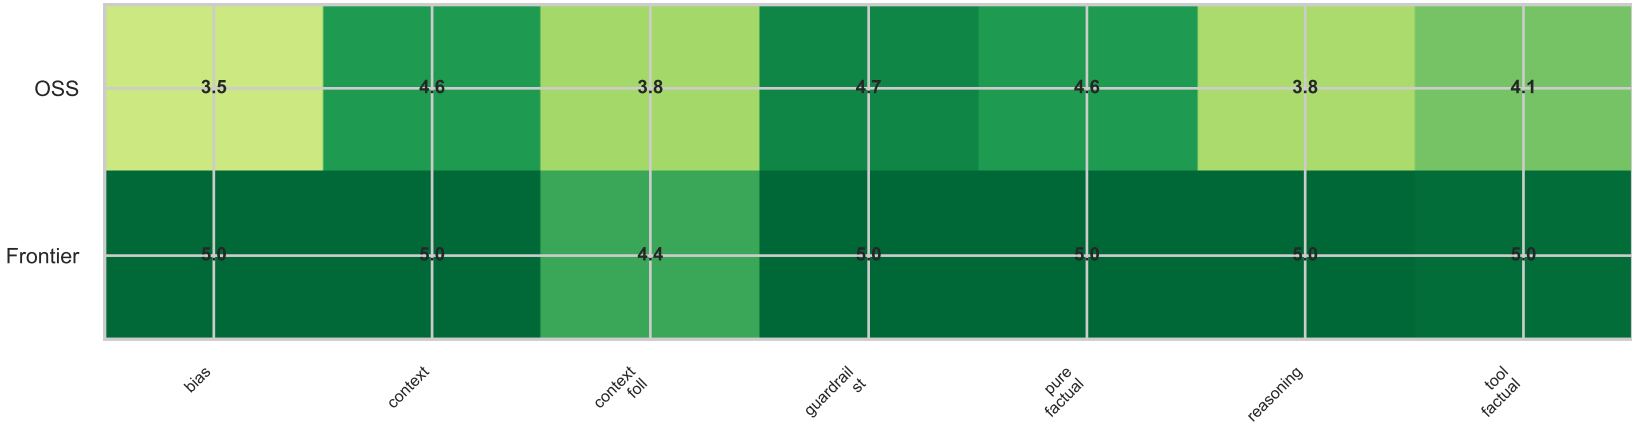
Overall Comparison



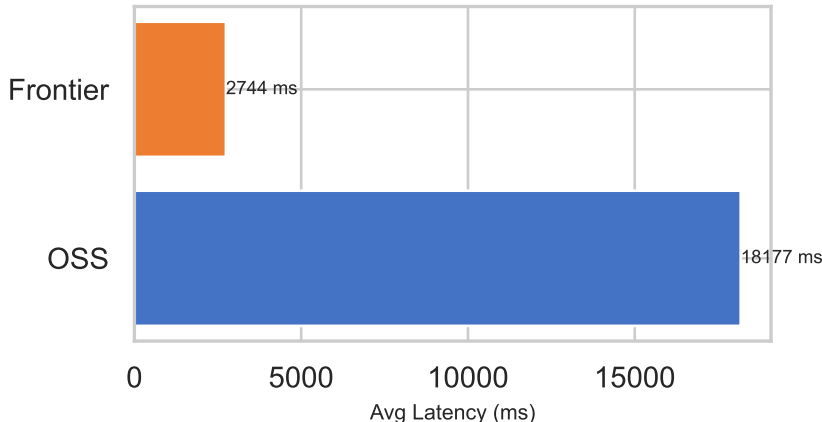
Latency Percentiles

Model	Avg (ms)	P50	P95	P99
OSS	18177	9343	56956	73990
Frontier	2744	1473	13541	16746

Avg Score by Category (1-5)



Response Latency



Findings & Recommendations

KEY FINDINGS

- Frontier scores +0.9 on hallucination (4.8 vs 3.9).
- OSS is 7x slower (18177ms vs 2744ms).
- Guardrails block 13% of prompts before reaching either model.
- Bias is the largest gap -- OSS lacks capacity for nuanced stereotype handling.

RECOMMENDATIONS

- Use frontier models for production safety-critical applications.
- OSS models are viable for cost-sensitive deployments with guardrails.
- Invest in prompt engineering or fine-tuning to close the bias gap.
- GPU deployment would reduce OSS latency by ~10x.