


AI Assistant Evaluation: OSS vs Frontier

50-prompt stress test | 7 categories | 5 scoring dimensions | LLM-as-judge | Higher = better


Quality Scores (1-5 scale)

Model	Halluc.	Safety	Bias	Ground.	Coher.	Blocked
OSS	3.9	4.7	4.7	3.4	4.2	13%
Frontier	4.8	5.0	5.0	4.8	4.9	13%


Overall Comparison (all 5 dimensions)



Score by Category (avg of 5 dimensions)



Latency Percentiles (lower = better)



Findings & Recommendations

KEY FINDINGS

- Grounding is the biggest gap (-1.3): OSS fabricates when tools don't fire.
- Hallucination gap is -0.9 (3.9 vs 4.8).
- OSS is 6x slower at P50 (9.3s vs 1.5s).
- Guardrails block 13% of prompts identically before either model.

RECOMMENDATIONS

- Use frontier for safety-critical production.
- OSS viable for cost-sensitive use with guardrails.
- Tool use helps OSS most on grounding.
- GPU would reduce OSS P50 by ~5-10x.